

Lab Assignment 2

Name: Soumyadeep Ganguly

Reg No.: 24MDT0082

```
In [256... import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
```

```
In [257... df = pd.read_csv('Training_set_heights200.csv')
df.head()
```

```
Out[257...      Height  Weight
0   127.8296   67.63371
1   123.4114   65.95421
2   134.4043   66.14316
3   155.9981   73.45251
4   136.1354   69.30943
```

```
In [258... from sklearn.preprocessing import MinMaxScaler
```

```
In [259... scaler = MinMaxScaler()
df_scaled = scaler.fit_transform(df)
```

```
In [260... X = df_scaled[:,0]
y = df_scaled[:,1]
```

```
In [261... from sklearn.linear_model import LinearRegression
```

```
In [262... LR = LinearRegression()
LR.fit(X.reshape(-1, 1), y)

print(f"Intercept: {LR.intercept_}")
print(f"Coefficient: {LR.coef_}")
```

Intercept: 0.2707514895889285
Coefficient: [0.52653218]

Splitting in Train and Test Set

```
In [263... from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_
```

```
In [264... LR2 = LinearRegression()
LR2.fit(X_train.reshape(-1, 1), y_train)
```

```
print(f"Intercept: {LR2.intercept_}")
print(f"Coefficient: {LR2.coef_}")
```

Intercept: 0.2612785266616628
Coefficient: [0.56093988]

Gradient Descent

```
In [265... def gd(xt, yt, m_b, n_b, L):
    D_m = 0
    D_n = 0
    for i in range(len(xt)):
        D_m = D_m + (2/len(xt))*((m_b*xt[i]+n_b - yt[i])*xt[i])
        D_n = D_n+(2/len(xt))*(m_b*xt[i]+n_b - yt[i])
    m_b = m_b - L*D_m
    n_b = n_b - L*D_n
    return m_b, n_b

m = 0
n = 0
L = 0.77
epochs = 900


for i in range(epochs):
    m,n = gd(X_train, y_train, m, n, L)

print(m,n)
```

0.5459409679484901 0.2340138965683014

```
In [266... plt.scatter(X_train, y_train)
y_pred = m*X_train+n
plt.plot(X_train, y_pred, color="red")
```

Out[266... [<matplotlib.lines.Line2D at 0x22e07a9fe00>]



```
In [267... y_test_pred = m*X_test + n
err = mean_squared_error(y_test, y_test_pred)
print(f"Testing Mean Squared Error: {err}")
```

Testing Mean Squared Error: 0.01576159799891117

Q1. Now create a linear regression model using gradient descent technique to predict height in terms of weight.

```
In [268... df2 = pd.read_csv('Training_set_heights200.csv')
df2.head()
```

```
Out[268...
   Height  Weight
0  127.8296  67.63371
1  123.4114  65.95421
2  134.4043  66.14316
3  155.9981  73.45251
4  136.1354  69.30943
```

```
In [269... scaler = MinMaxScaler()
df_scaled = scaler.fit_transform(df)
```

height in terms of weight

```
In [270... X = df_scaled[:,1]
y = df_scaled[:,0]
```

```
In [271... from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_
```

```
In [272... LR3 = LinearRegression()
LR3.fit(X_train.reshape(-1, 1), y_train)

print(f"Intercept: {LR3.intercept_}")
print(f"Coefficient: {LR3.coef_}")
```

Intercept: 0.24649525701785846
Coefficient: [0.51939401]

Iteration till 4 decimal points unchanged

```
In [273... def gradient_decent(xt, yt, m_b, n_b, L):
    D_m = 0
    D_n = 0
    for i in range(len(xt)):
        D_m += (2 / len(xt)) * ((m_b * xt[i] + n_b - yt[i]) * xt[i])
        D_n += (2 / len(xt)) * (m_b * xt[i] + n_b - yt[i])
    m_b_new = m_b - L * D_m
    n_b_new = n_b - L * D_n
    return m_b_new, n_b_new

def run_gradient_descent(xt, yt, m_b, n_b, L, max_iterations=1000):
    for iteration in range(max_iterations):
        m_b_new, n_b_new = gradient_decent(xt, yt, m_b, n_b, L)
        if np.all(np.round(m_b_new, 4) == np.round(m_b, 4)) and np.all(np.round(
            print(f"Converged after {iteration + 1} iterations.")
            break
        m_b, n_b = m_b_new, n_b_new
    else:
        print("Maximum iterations reached without full convergence.")
    return m_b, n_b

m_b = 0.0
n_b = 0.0
learning_rate = 0.2


m_b, n_b = run_gradient_descent(X_train, y_train, m_b, n_b, learning_rate)
print(f"Final parameters: m_b = {m_b}, n_b = {n_b}")
```

Converged after 393 iterations.

Final parameters: m_b = 0.5085540073514879, n_b = 0.2527363774675484

```
In [274... plt.scatter(X_train, y_train)
y_pred = m_b*X_train+n_b
plt.plot(X_train, y_pred, color="red")
```

```
Out[274... [<matplotlib.lines.Line2D at 0x22e09566b10>]
```



```
In [275... y_test_pred = m*X_test + n
err = mean_squared_error(y_test, y_test_pred)
print(f"Testing Mean Squared Error: {err}")
```

Testing Mean Squared Error: 0.018597048067984106

Multiple Linear Regression

Application on Book1.csv file

```
In [276... book = pd.read_csv('Book1.csv')
book.head()
```

```
Out[276...
   price  area  bedrooms  bathrooms  stories  parking  furnishingstatus
0  13300000  7420         4          2        3         2      furnished
1  12250000  8960         4          4        4         3      furnished
2  12250000  9960         3          2        2         2  semi-furnished
3  12215000  7500         4          2        2         3      furnished
4  11410000  7420         4          1        2         2      furnished
```

```
In [277... book.shape
```

```
Out[277... (249, 7)
```

```
In [278... book.isna().sum()
```

```
Out[278... price          0
          area          0
          bedrooms      0
          bathrooms     0
          stories       0
          parking       0
          furnishingstatus 0
          dtype: int64
```

```
In [279... book['furnishingstatus'].unique()
```

```
Out[279... array(['furnished', 'semi-furnished', 'unfurnished'], dtype=object)
```

```
In [280... X_train, X_test, y_train, y_test = train_test_split(book.drop(['furnishingstatus
```

1-hot encoding

```
In [281... from sklearn.preprocessing import OneHotEncoder
```

```
In [282... encoder = OneHotEncoder(sparse_output=False)
one_hot_encoded = encoder.fit_transform(book[['furnishingstatus']])
```

```
In [283... one_hot_book = pd.DataFrame(one_hot_encoded, columns=encoder.get_feature_names_o
```

```
In [284... book_encoded = pd.concat([book, one_hot_book], axis=1)
book_encoded.drop(['furnishingstatus'], axis=1, inplace=True)
book_encoded
```

```
Out[284...           price  area  bedrooms  bathrooms  stories  parking  furnishingstatus_furnisher
```

	price	area	bedrooms	bathrooms	stories	parking	furnishingstatus_furnisher
0	13300000	7420	4	2	3	2	1.
1	12250000	8960	4	4	4	3	1.
2	12250000	9960	3	2	2	2	0.
3	12215000	7500	4	2	2	3	1.
4	11410000	7420	4	1	2	2	1.
...
244	4550000	5320	3	1	2	0	0.
245	4550000	5360	3	1	2	2	0.
246	4550000	3520	3	1	1	0	0.
247	4550000	8400	4	1	4	3	0.
248	4543000	4100	2	2	1	0	0.

249 rows × 9 columns



```
In [285... X = book_encoded.drop(['price'], axis=1)
y = book_encoded['price']
```

```
y = np.array(y).reshape(-1, 1)
```

Preprocessing

```
In [286... from sklearn.preprocessing import StandardScaler
```

```
SS = StandardScaler()
X_scaled = SS.fit_transform(X)
y_scaled = SS.fit_transform(y)
```

```
In [287... X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_scaled, test_siz
```

```
In [288... y_train = np.array(y_train.reshape(199))
```

```
y_train.shape
```

```
Out[288... (199,)
```

```
In [289... m_b = 0.0
n_b = 0.0
learning_rate = 0.2

m_b, n_b = run_gradient_descent(X_train, y_train, m_b, n_b, learning_rate)
print(f"Final parameters: m_b = {m_b}, n_b = {n_b}")
```

Converged after 19 iterations.

```
Final parameters: m_b = [ 0.28879476  0.30878645  0.46680269  0.25548182  0.30603
175  0.27005068
-0.20558665 -0.05654349], n_b = [-0.01839882 -0.01134067 -0.01044147 -0.03059856
-0.00618137 -0.01388126
-0.01126489 -0.02028817]
```

```
Final parameters: m_b = [ 0.28879476  0.30878645  0.46680269  0.25548182  0.30603
175  0.27005068
-0.20558665 -0.05654349], n_b = [-0.01839882 -0.01134067 -0.01044147 -0.03059856
-0.00618137 -0.01388126
-0.01126489 -0.02028817]
```

```
In [290... y_p = 0
for m in m_b:
    y_p = y_p + m*X_test

y_p = y_p + n_b[-1]
```

```
In [291... y_p.shape
```

```
Out[291... (50, 8)
```

```
In [ ]:
```

```
In [ ]:
```