

ML LAB 6

Name: Soumyadeep Ganguly

Reg No: 24MDT0082

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv("liver_patient.csv")
df.head()
```

```
Out[2]:
```

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Proti
0	65	Female	0.7	0.1	187	16	18	
1	62	Male	10.9	5.5	699	64	100	
2	62	Male	7.3	4.1	490	60	68	
3	58	Male	1.0	0.4	182	14	20	
4	72	Male	3.9	2.0	195	27	59	



```
In [3]: df.drop(['Age', 'Gender'], axis=1, inplace=True)
```


```
In [4]: from sklearn.preprocessing import MinMaxScaler

mms = MinMaxScaler(feature_range=(0,1))
data = mms.fit_transform(df)
cols = df.columns[:]
```

```
df = pd.DataFrame(data=data, columns=cols)
df.head(4)
```

Out[4]:

	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Protiens	Albumin
0	0.004021	0.000000	0.060576	0.003015	0.001626	0.594203	0.521739
1	0.140751	0.275510	0.310699	0.027136	0.018296	0.695652	0.500000
2	0.092493	0.204082	0.208598	0.025126	0.011791	0.623188	0.521739
3	0.008043	0.015306	0.058134	0.002010	0.002033	0.594203	0.543478



In [5]: d = df.values
d

Out[5]: array([[0.00402145, 0. , 0.06057645, ..., 0.52173913, 0.24 ,
1.],
[0.14075067, 0.2755102 , 0.31069858, ..., 0.5 , 0.176 ,
1.],
[0.0924933 , 0.20408163, 0.20859795, ..., 0.52173913, 0.236 ,
1.],
...,
[0.00536193, 0.00510204, 0.0889106 , ..., 0.5 , 0.28 ,
1.],
[0.01206434, 0.02040816, 0.05911089, ..., 0.54347826, 0.28 ,
1.],
[0.0080429 , 0.01020408, 0.07474353, ..., 0.76086957, 0.48 ,
0.]], shape=(583, 9))

In [6]: x = d[:, :-1]
y = d[:, -1]

In [7]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=0)

In [9]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

```
l_reg = LogisticRegression()  
l_reg.fit(x_train,y_train)
```

Out[9]:

```
▼ LogisticRegression ⓘ ?  
LogisticRegression()
```

```
In [10]: y_pred = l_reg.predict(x_train)  
acc_lr = accuracy_score(y_pred, y_train)  
acc_lr
```

Out[10]: 0.7253218884120172

K-Fold Cross Validation

```
In [11]: from sklearn.model_selection import KFold, cross_val_score  
  
k_f_val = KFold(n_splits=5, shuffle=True, random_state=0)  
results = cross_val_score(l_reg, x_train, y_train, scoring='accuracy', cv=k_f_val)  
results
```

Out[11]: array([0.81914894, 0.72043011, 0.74193548, 0.70967742, 0.6344086])

```
In [12]: np.mean(results)
```

Out[12]: np.float64(0.7251201098146878)

Stratified K-Fold Cross Validation

```
In [13]: from sklearn.model_selection import StratifiedKFold  
  
skf_val = StratifiedKFold(n_splits=5, shuffle=True, random_state=0)  
  
results_skf = cross_val_score(l_reg, x_train, y_train, scoring='accuracy', cv=skf_val)  
results_skf
```

```
Out[13]: array([0.72340426, 0.72043011, 0.72043011, 0.7311828 , 0.7311828 ])
```

```
In [14]: results_skf.mean()
```

```
Out[14]: np.float64(0.7253260123541525)
```

Performance Measures

```
In [15]: y_pred_test = l_reg.predict(x_test)
accuracy_score(y_pred_test, y_test)
```

```
Out[15]: 0.6666666666666666
```

```
In [16]: from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score, ConfusionMatrixDisplay
cm = confusion_matrix(y_test, y_pred_test)
cm
```

```
Out[16]: array([[ 0, 39],
               [ 0, 78]])
```

```
In [17]: ps = precision_score(y_test, y_pred_test)
ps
```

```
Out[17]: 0.6666666666666666
```

```
In [18]: rs = recall_score(y_test, y_pred_test)
rs
```

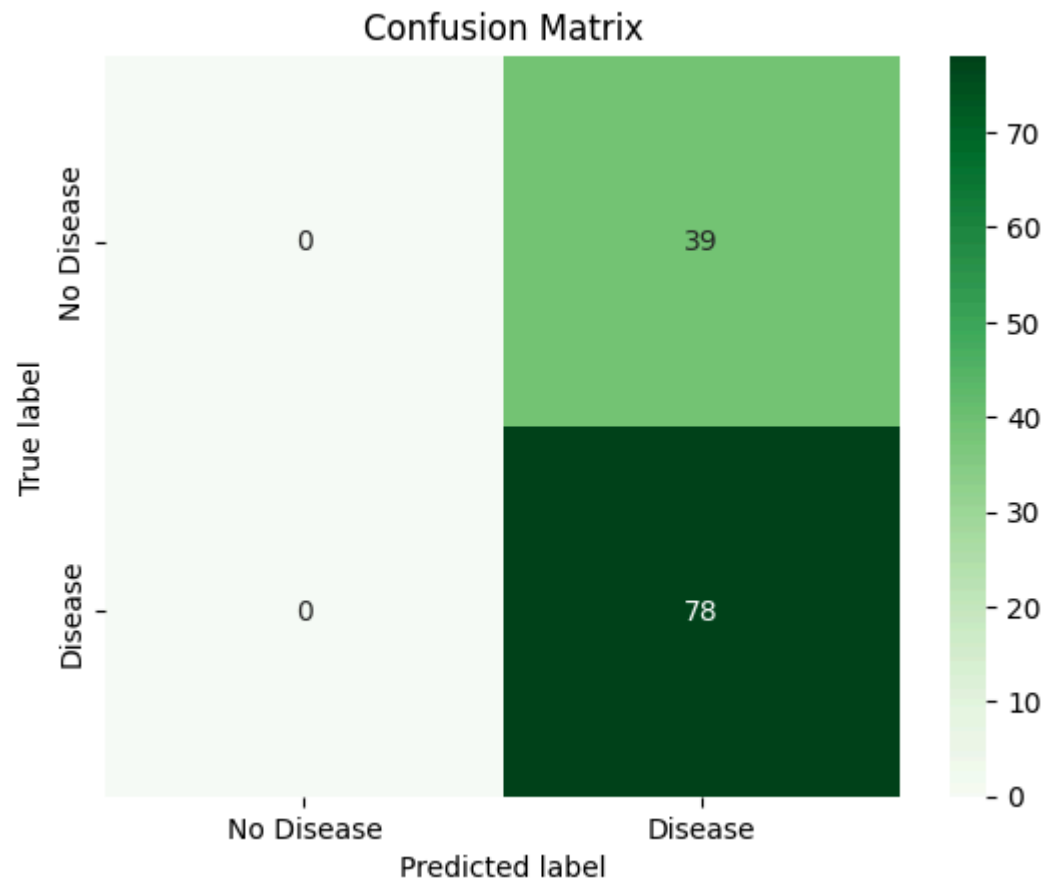
```
Out[18]: 1.0
```

```
In [19]: f1s = f1_score(y_test, y_pred_test)
f1s
```

```
Out[19]: 0.8
```

```
In [20]: import seaborn as sns
sns.heatmap(cm, annot=True, fmt='d', cmap='Greens', xticklabels=['No Disease', 'Disease'], yticklabels=['No Disease', 'Disease'])
```

```
plt.xlabel('Predicted label')  
plt.ylabel('True label')  
plt.title('Confusion Matrix')  
plt.show()
```



Decision Tree - logistic Regression

```
In [21]: from sklearn.tree import DecisionTreeClassifier, plot_tree  
dtc = DecisionTreeClassifier()  
dtc.fit(x_train, y_train)
```

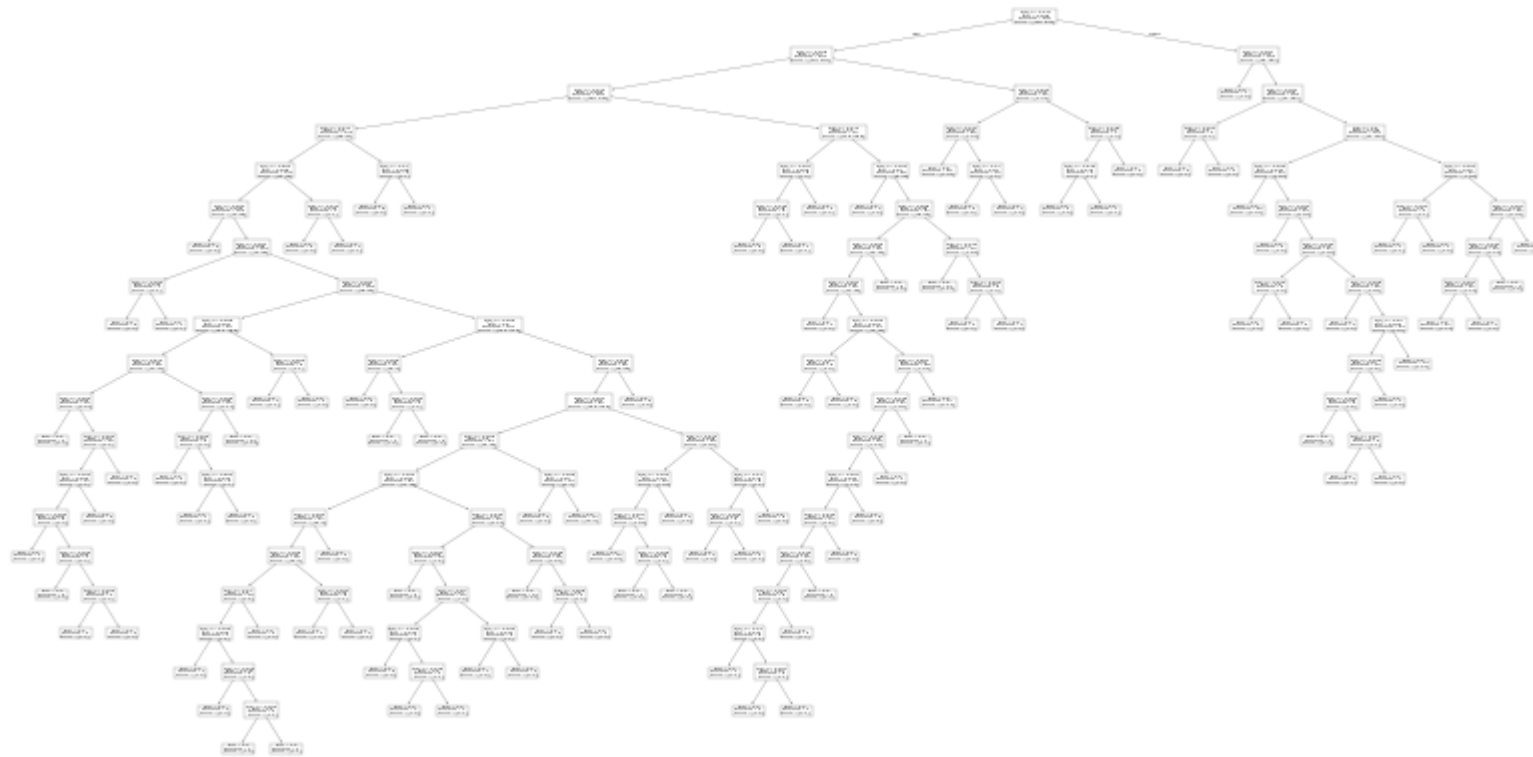
Out[21]: ▾ DecisionTreeClassifier ⓘ ?

DecisionTreeClassifier()

In [23]: `pred = dtc.predict(x_test)`
`accuracy_score(pred, y_test)`

Out[23]: 0.6495726495726496

In [24]: `plt.figure(figsize=(100,50), dpi=10)`
`plot_tree(dtc)`
`plt.show()`



Decision Tree - Linear Regression

```
In [25]: df2 = pd.read_csv("Book1.csv")
df2.head()
```

```
Out[25]:
```

	price	area	bedrooms	bathrooms	stories	parking	furnishingstatus
0	13300000	7420	4	2	3	2	furnished
1	12250000	8960	4	4	4	3	furnished
2	12250000	9960	3	2	2	2	semi-furnished
3	12215000	7500	4	2	2	3	furnished
4	11410000	7420	4	1	2	2	furnished

```
In [26]: df2.drop(['furnishingstatus'], axis=1, inplace=True)
```

```
In [27]: mms = MinMaxScaler(feature_range=(0,1))
d = mms.fit_transform(df2)
cols = df2.columns[:]
df2 = pd.DataFrame(data=d, columns=cols)
df2.head(4)
```

```
Out[27]:
```

	price	area	bedrooms	bathrooms	stories	parking
0	1.000000	0.356777	0.50	0.333333	0.666667	0.666667
1	0.880096	0.469597	0.50	1.000000	1.000000	1.000000
2	0.880096	0.542857	0.25	0.333333	0.333333	0.666667
3	0.876099	0.362637	0.50	0.333333	0.333333	1.000000

```
In [28]: x = df2.values[:,1:]  
y = df2.values[:,1]  
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=0)
```

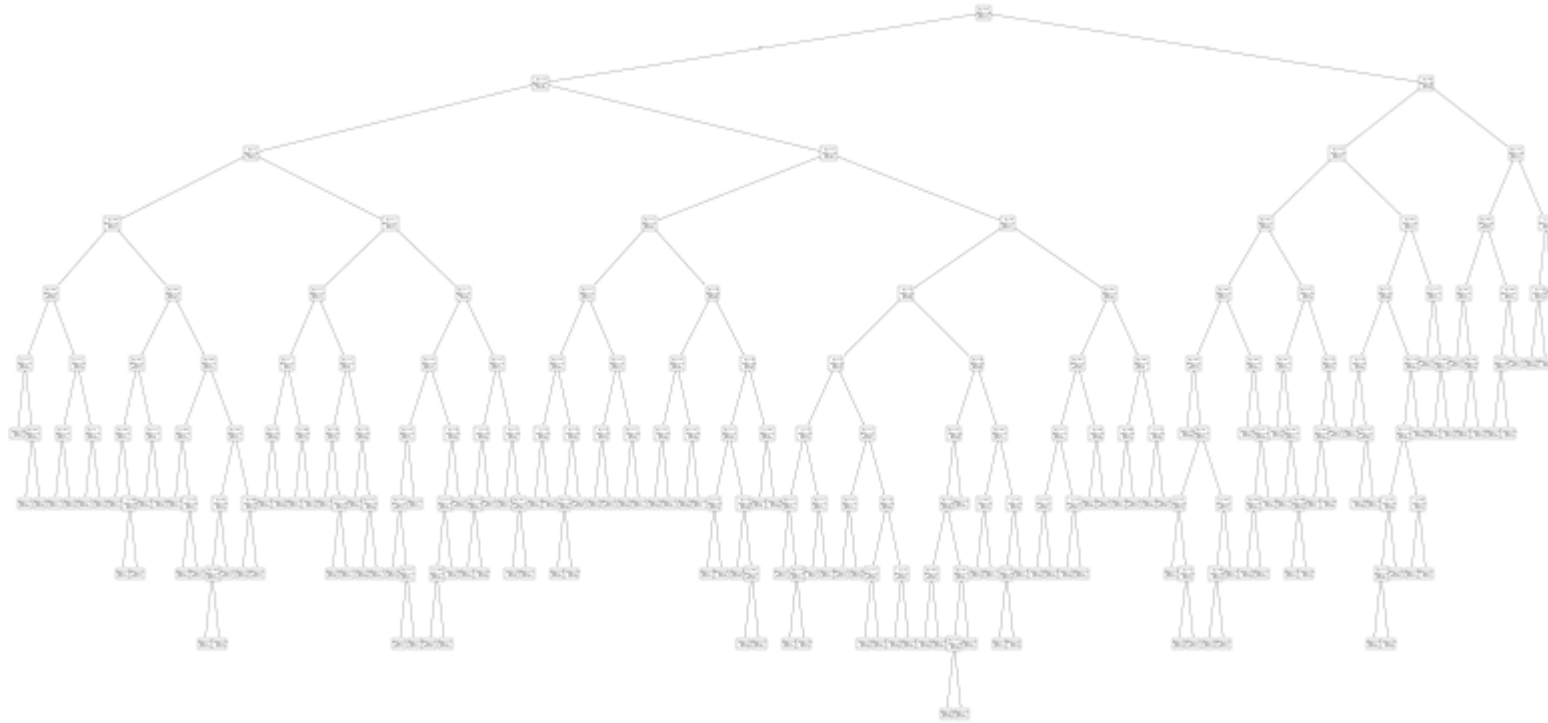
```
In [29]: from sklearn.tree import DecisionTreeRegressor  
from sklearn.metrics import mean_squared_error  
  
dtt_linreg = DecisionTreeRegressor()  
dtt_linreg.fit(x_train, y_train)
```

```
Out[29]: ▾ DecisionTreeRegressor ⓘ ?  
DecisionTreeRegressor()
```

```
In [30]: pred = dtt_linreg.predict(x_test)  
mean_squared_error(pred, y_test)
```

```
Out[30]: 0.0006450250506514239
```

```
In [31]: plt.figure(figsize=(100,50), dpi=10)  
plot_tree(dtt_linreg)  
plt.show()
```

Hyperparameter Tuning: Grid Search

```
In [32]: from sklearn.model_selection import GridSearchCV
```

```
In [33]: df = pd.read_csv('liver_patient.csv')  
df.head(4)
```

Out[33]:

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferase	Total_Proti
0	65	Female	0.7	0.1	187	16	18	
1	62	Male	10.9	5.5	699	64	100	
2	62	Male	7.3	4.1	490	60	68	
3	58	Male	1.0	0.4	182	14	20	

In [34]: `df.drop('Gender', axis=1, inplace=True)`

In [35]:

```

mms = MinMaxScaler()
x = mms.fit_transform(df)
X = x[:, 0:9]
Y = x[:, 9]
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

```

In [36]:

```

param_grid = {
    'max_depth': [3, 5, 10],
    'min_samples_leaf': [1, 5, 10, 20]
}

```

In [37]:

```

dt_model = DecisionTreeClassifier(random_state=0)
grid_search = GridSearchCV(estimator=dt_model, param_grid=param_grid, cv=5, scoring='accuracy', n_jobs=-1)
grid_search.fit(X_train, Y_train)
best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
Y_pred = best_model.predict(X_test)
accuracy = accuracy_score(Y_test, Y_pred)
print("Best Parameters:", best_params)
print("Best Cross-validation Accuracy:", grid_search.best_score_)
print("Test Accuracy:", accuracy)

```

Best Parameters: {'max_depth': 3, 'min_samples_leaf': 1}
 Best Cross-validation Accuracy: 0.6759780370624571
 Test Accuracy: 0.717948717948718

In []: