

ML Lab 10: 26th March, 2025

Name: Soumyadeep Ganguly

Reg No: 24MDT0082

```
In [62]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [63]: df = pd.read_csv('lab10data.csv')
df.head()
```

```
Out[63]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	48	1	2	124	255	1	1	175	0	0.0	2	2	2	1
1	68	0	2	120	211	0	0	115	0	1.5	1	0	2	1
2	46	1	0	120	249	0	0	144	0	0.8	2	0	3	0
3	60	1	0	130	253	0	1	144	1	1.4	2	1	3	0
4	43	1	0	115	303	0	1	181	0	1.2	1	0	2	1

```
In [64]: df.dtypes
```

```
Out[64]: age      int64
sex      int64
cp       int64
trestbps int64
chol     int64
fbs      int64
restecg  int64
thalach  int64
exang    int64
oldpeak  float64
slope    int64
ca       int64
thal     int64
target   int64
dtype: object
```

Preprocessing

```
In [65]: from sklearn.preprocessing import MinMaxScaler
```

```
In [66]: scaler = MinMaxScaler()
X = df.drop(['target'], axis=1)
y = df['target']

scaled_X = scaler.fit_transform(X)
X = pd.DataFrame(scaled_X, columns = X.columns)
X.head()
```

```
Out[66]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	0.395833	1.0	0.666667	0.306122	0.294521	1.0	0.5	0.793893	0.0	0.000000	1.0	0.50	0.666667
1	0.812500	0.0	0.666667	0.265306	0.194064	0.0	0.0	0.335878	0.0	0.267857	0.5	0.00	0.666667
2	0.354167	1.0	0.000000	0.265306	0.280822	0.0	0.0	0.557252	0.0	0.142857	1.0	0.00	1.000000
3	0.645833	1.0	0.000000	0.367347	0.289954	0.0	0.5	0.557252	1.0	0.250000	1.0	0.25	1.000000
4	0.291667	1.0	0.000000	0.214286	0.404110	0.0	0.5	0.839695	0.0	0.214286	0.5	0.00	0.666667

PCA

```
In [67]: from sklearn.decomposition import PCA
```

```
In [68]: pca = PCA(n_components = 2)
pc_X = pca.fit_transform(X)
pc_X
```

```
Out[68]: array([[-0.23522174, -0.7111645 ],
               [-0.61442471,  0.48110119],
               [-0.06443809, -0.35336159],
               [ 0.69975976,  0.08597526],
               [-0.05093584, -0.30573675],
               [ 0.1003509 , -0.2138113 ],
               [ 0.74877193,  0.09960479],
               [-0.21849872, -0.68918824],
               [ 0.46663659, -0.013086  ],
               [-0.62516878,  0.48582954],
               [-0.624672 ,  0.44984363],
               [-0.36929451, -0.53567358],
               [ 0.78364549, -0.12752742],
               [-0.81717358,  0.34101822],
               [ 0.235398 ,  1.00066192],
               [-0.29051287,  0.59591091],
               [ 0.71093233,  0.12279361],
               [ 0.86473216,  0.17145438],
               [-0.20406023, -0.43625508],
               [ 0.61705649, -0.14238798],
               [-0.62827583,  0.45517818],
               [ 0.83791947,  0.22383137],
               [ 0.69987153,  0.07697176],
               [-0.66440001,  0.43927683],
               [-0.47643074,  0.61143065],
               [-0.1259226 , -0.43429024],
               [ 0.07064487, -0.549728  ],
               [-0.83387913,  0.34080572],
               [-0.03930769, -0.36759135],
               [ 0.26501976, -0.41687917],
               [-0.23167894, -0.68999279],
               [ 0.8776946 ,  0.18705249],
               [ 0.11930381, -0.23791664],
               [ 0.31762362, -0.21318989],
               [ 0.59730846,  0.03852486],
               [-0.08579792, -0.36699754],
               [-0.30030033, -0.51949687],
               [-0.85538856,  0.39287624],
               [-0.7251578 ,  0.38146  ],
               [-0.10311769, -0.36264012],
               [ 0.96964597, -0.07519052],
               [ 0.71831183,  0.18420095],
               [ 0.27177819,  0.98294635],
               [ 0.00273046, -0.28721551],
               [-0.78589445,  0.39571354],
```

[-0.27394347, -0.44914303],
[0.52823321, -0.0065165],
[-0.12888377, -0.34046886],
[0.70976969, 0.10056697],
[0.71949439, 0.18620373],
[0.85154601, 0.10942926],
[0.84396617, 0.19642803],
[0.91642401, 0.21855696],
[0.0298716 , -0.30259733],
[0.63508579, 0.03703488],
[0.17108728, -0.21097597],
[0.42856007, 0.8406353],
[-0.07481201, -0.26284782],
[0.35098562, -0.12584041],
[0.01194653, -0.32054236],
[-0.82226992, 0.31321272],
[-0.7013165 , 0.21640359],
[-0.67365892, 0.47763115],
[-0.07954792, -0.31912431],
[-0.20317366, -0.37465955],
[-0.39219682, -0.55098082],
[0.8348998 , 0.1511027],
[-0.27350548, -0.47464119],
[-0.81183902, 0.29909735],
[-0.13529798, -0.39928764],
[-0.63536939, 0.50345193],
[0.88389682, 0.20636924],
[0.10874391, -0.27944135],
[-0.05540845, -0.35300344],
[0.82227196, 0.18383936],
[-0.24595949, 0.45774802],
[0.85361623, 0.18058447],
[-0.19231146, -0.4148893],
[0.83233772, 0.15275175],
[-0.04685948, -0.59919274],
[-0.69689453, 0.40800494],
[-0.82288329, 0.34240302],
[-0.49051069, 0.57321495],
[-0.78690786, 0.33419871],
[-0.1699777 , -0.38990308],
[-0.2496717 , -0.5849888],
[0.91334767, 0.1698216],
[-0.69708937, 0.38014988],
[0.48964409, -0.11973564],
[-0.72799235, 0.42992157],

[-0.4500464 , 0.61183695],
[-0.37848341, -0.53904371],
[0.62010297, 0.02533211],
[-0.15672572, -0.35441692],
[0.69133623, -0.13072716],
[0.30166293, 1.05769685],
[-0.03121024, -0.36680337],
[0.24468684, -0.13949864],
[0.77177038, 0.20591734],
[-0.14727578, -0.49224378],
[0.89374386, 0.19242046],
[-0.17158696, -0.38913822],
[-0.88820128, 0.26698651],
[-0.22046594, -0.4320584],
[0.38173298, 1.00707193],
[-0.7009914 , 0.38214912],
[-0.73316371, 0.36828529],
[-0.5704898 , 0.25681803],
[0.11323205, 0.65773631],
[0.90668187, 0.19946893],
[-0.2530694 , -0.50313929],
[-0.93432507, 0.26490167],
[0.87200298, 0.15864359],
[0.26951735, 0.98982208],
[0.00190079, -0.39316829],
[0.08328001, -0.17026781],
[0.94134428, 0.15982886],
[-0.12712544, -0.38748441],
[-0.38690841, -0.53663255],
[-0.17486367, -0.50418058],
[0.19363302, -0.36918903],
[-0.36682692, -0.52585768],
[-0.13137183, -0.37031337],
[-0.80370545, 0.34505265],
[0.66865293, 0.10271856],
[-0.2920301 , -0.49137884],
[0.40654496, -0.12848371],
[-0.25579969, -0.42174571],
[-0.71229089, 0.3821404],
[-0.03629649, -0.42102783],
[-0.80931678, 0.37575319],
[1.03983806, 0.10360788],
[-0.08074347, -0.54354588],
[-0.25795541, -0.43661242],
[0.05576172, -0.37762175],

[-0.63326199, 0.47582315],
[-0.08702181, -0.53985418],
[-0.35850425, -0.53209181],
[0.70922636, -0.01000493],
[-0.77310781, 0.37297937],
[-0.59725764, 0.46102883],
[-0.14773805, -0.36158498],
[0.56966999, -0.01330883],
[0.68424026, 0.09443224],
[-0.47121715, 0.57057377],
[-0.28327865, -0.50005278],
[-0.32721012, -0.50450635],
[-0.08179176, -0.36726268],
[-0.59316511, 0.49405509],
[-0.25153947, -0.48766431],
[-0.8788213 , 0.27719382],
[-0.14809461, -0.63563678],
[-0.25469626, -0.497448],
[0.83762707, 0.1150745],
[-0.10650819, -0.40057984],
[-0.19884895, -0.46223222],
[1.02358416, 0.05598585],
[-0.16826454, -0.35682576],
[-0.49441659, 0.57858832],
[-0.74128377, 0.3517146],
[-0.21490447, -0.71525618],
[0.26127303, -0.40031193],
[-0.23975655, -0.47321503],
[-0.18483861, -0.4454762],
[0.14439287, -0.4547836],
[0.29677479, 1.00489791],
[-0.66428874, 0.43257768],
[-0.08249778, -0.38468823],
[-0.14012566, -0.40383759],
[-0.25156355, 0.62408792],
[0.20641428, 0.99096865],
[-0.08351051, -0.33645949],
[-0.16785057, -0.69981446],
[0.83747295, 0.14523762],
[1.03933157, 0.29595457],
[-0.2496717 , -0.5849888],
[-0.25333013, -0.44098314],
[0.801883 , 0.14942923],
[0.97604494, 0.18866931],
[-0.27176885, -0.43117751],

```
[-0.76954875, 0.38330075],  
[ 0.89997024, -0.05208366],  
[ 0.08661972, 0.92021675],  
[ 0.80555159, 0.1680903 ],  
[-0.12288882, 0.71289617],  
[-0.6557288 , 0.42744003],  
[-0.37212588, 0.62258391],  
[ 0.6183415 , -0.14845285],  
[-0.1579606 , -0.41669491],  
[ 0.79424775, -0.10816758],  
[-0.26497277, -0.45423412],  
[ 0.69625181, 0.11059579],  
[-0.76758338, 0.35077322],  
[-0.33058104, 0.57816422],  
[ 0.90997533, 0.19102861],  
[ 0.21534599, 0.99049656],  
[-0.63087125, 0.48060299],  
[ 0.98804415, 0.30064055],  
[-0.20202485, -0.42794201],  
[-0.1404868 , -0.36582407],  
[-0.27134433, -0.49572141],  
[-0.67002997, 0.22076496],  
[-0.16163726, -0.62289309],  
[-0.56186612, 0.46691356],  
[-0.12862635, -0.56504381],  
[-0.26096797, -0.44002333],  
[-0.25734822, -0.66957447],  
[ 0.92717704, 0.19158558],  
[-0.81805936, 0.11660116],  
[ 0.91401722, 0.18041303],  
[-0.65465385, 0.48295743],  
[ 0.41717656, 0.85663035]])
```

(a) Logistic Regression

```
In [69]: from sklearn.linear_model import LogisticRegression  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score, confusion_matrix
```


```
In [70]: X_train, X_test, y_train, y_test = train_test_split(pc_X, y, test_size = 0.2, random_state = 42)
```



```
In [71]: LR = LogisticRegression()
LR.fit(X_train, y_train)
lr_y_pred = LR.predict(X_test)
lr_cfm = confusion_matrix(y_test, lr_y_pred)
print(f"Accuracy of Logistic Regression is: {accuracy_score(y_test, lr_y_pred)}")
```

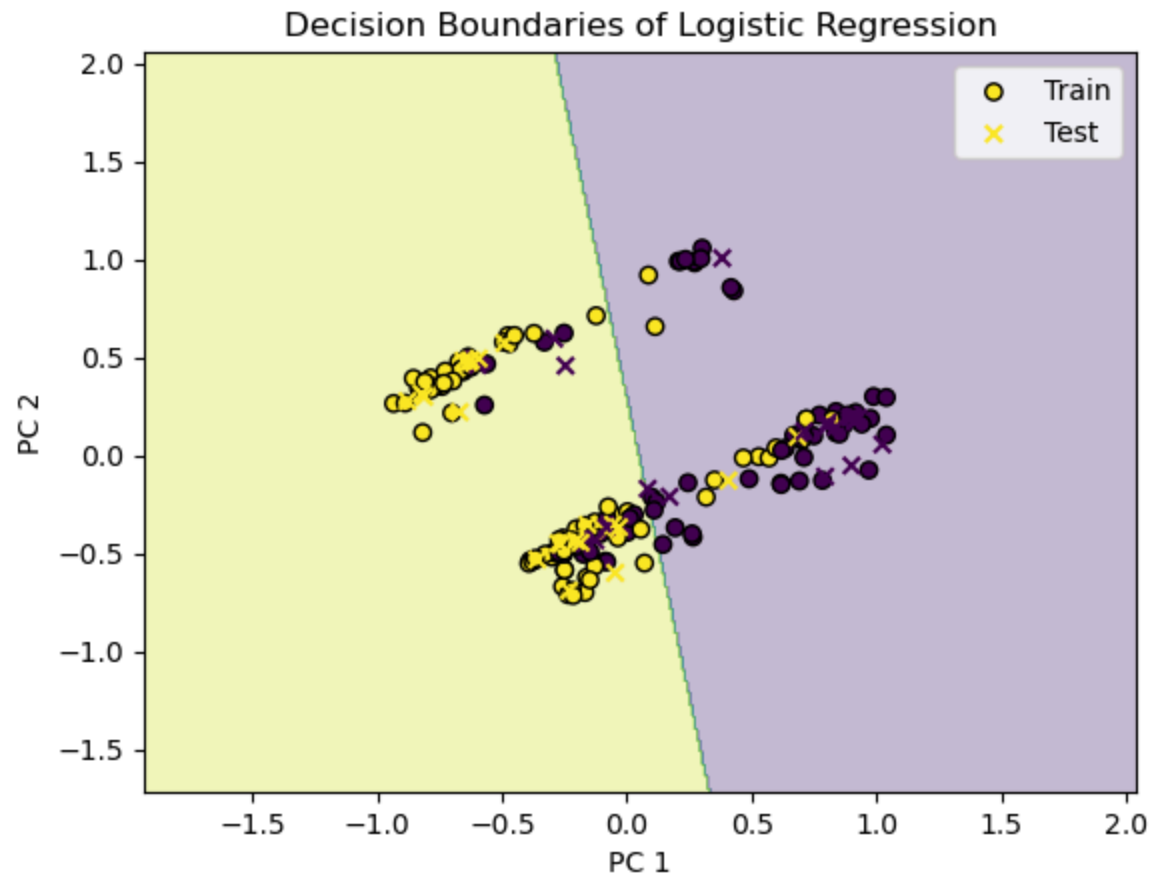
Accuracy of Logistic Regression is: 0.8372093023255814

```
In [72]: sns.heatmap(lr_cfm, annot=True)
plt.title('Confusion Matrix for Logistic Regression')
plt.show()
```



```
In [73]: x_min, x_max = pc_X[:, 0].min() -1, pc_X[:, 0].max() +1
y_min, y_max = pc_X[:, 1].min() -1, pc_X[:, 1].max() +1
xx, yy = np.meshgrid(np.linspace(x_min, x_max, 400), np.linspace(y_min, y_max, 400))
Z = LR.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)
```

```
plt.contourf(xx, yy, Z, alpha=0.3)
plt.scatter(X_train[:, 0], X_train[:, 1], c=y_train, edgecolors='k', label='Train')
plt.scatter(X_test[:, 0], X_test[:, 1], c=y_test, marker='x', label='Test')
plt.xlabel("PC 1")
plt.ylabel("PC 2")
plt.title("Decision Boundaries of Logistic Regression")
plt.legend()
plt.show()
```




(b) Decision tree

```
In [74]: from sklearn.tree import DecisionTreeClassifier, plot_tree
```

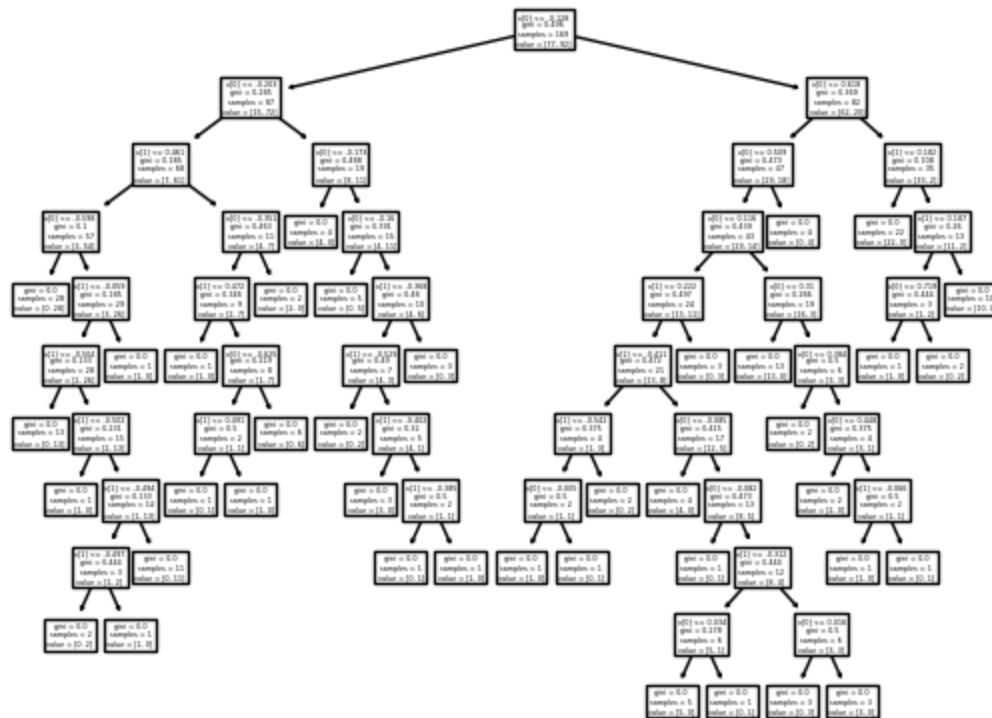
```
In [75]: DT = DecisionTreeClassifier()
DT.fit(X_train, y_train)
dt_y_pred = DT.predict(X_test)
dt_cfm = confusion_matrix(y_test, dt_y_pred)
print(f"Accuracy of Decision Tree Classifier is: {accuracy_score(y_test, dt_y_pred)}")
```

Accuracy of Decision Tree Classifier is: 0.7906976744186046

```
In [76]: sns.heatmap(dt_cfm, annot=True)
plt.title('Confusion Matrix for Decision Tree Classifier')
plt.show()
```



```
In [77]: plot_tree(DT)
plt.show()
```




(c) Bagging

```
In [78]: from sklearn.ensemble import BaggingClassifier
BC = BaggingClassifier(n_estimators = 100, random_state = 0)
BC.fit(X_train,y_train)
bc_y_pred = BC.predict(X_test)

dt_cfm = confusion_matrix(y_test, bc_y_pred)
print(f"Accuracy of Decision Tree Classifier is: {accuracy_score(y_test, bc_y_pred)}")

sns.heatmap(dt_cfm, annot=True)
plt.title('Confusion Matrix for Decision Tree Classifier')
plt.show()
```

Accuracy of Decision Tree Classifier is: 0.7906976744186046



Hyperparameter Tuning on Bagging Classifier

```
In [79]: from sklearn.model_selection import GridSearchCV
bc_param_grid = {
    'n_estimators': [10, 50, 80, 100, 150, 200, 250, 300]
}
bc_base_model = BaggingClassifier(random_state=0)

bc_gs = GridSearchCV(estimator=bc_base_model, param_grid=bc_param_grid, cv=5, n_jobs=-1)
bc_gs.fit(X_train, y_train)

print(f'Best parameters for "Bagging Classifier": {bc_gs.best_params_}')
print(f'Best Accuracy: {bc_gs.best_score_}')
```

Best parameters for "Bagging Classifier": {'n_estimators': 50}
 Best Accuracy: 0.7392156862745097

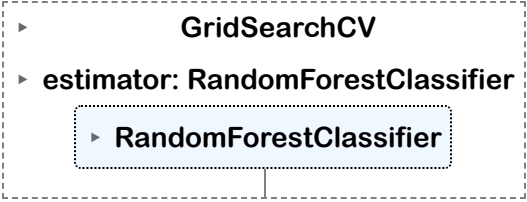
(d) RandomForest with hyperparamter tuning

```
In [80]: from sklearn.ensemble import RandomForestClassifier
```

```
In [81]: rfc_base_model = RandomForestClassifier(max_features='sqrt')
rfc_param_grid = {
    'n_estimators': [10, 50, 80, 100, 150, 200, 250, 300],
    'max_depth': [1, 2, 3, 4, 10, 15, 20]
}

rfc_gs = GridSearchCV(estimator = rfc_base_model, param_grid = rfc_param_grid, cv=5, n_jobs = -1)
rfc_gs.fit(X_train, y_train)
```

```
Out[81]:
```



```

  ▸ GridSearchCV
  ▸ estimator: RandomForestClassifier
      ▸ RandomForestClassifier

```

```
In [82]: print(f'Best parameters for "Random Forest Classifier": {rfc_gs.best_params_}')
print(f'Best Accuracy: {rfc_gs.best_score_}')
```

Best parameters for "Random Forest Classifier": {'max_depth': 2, 'n_estimators': 50}
 Best Accuracy: 0.7811051693404634

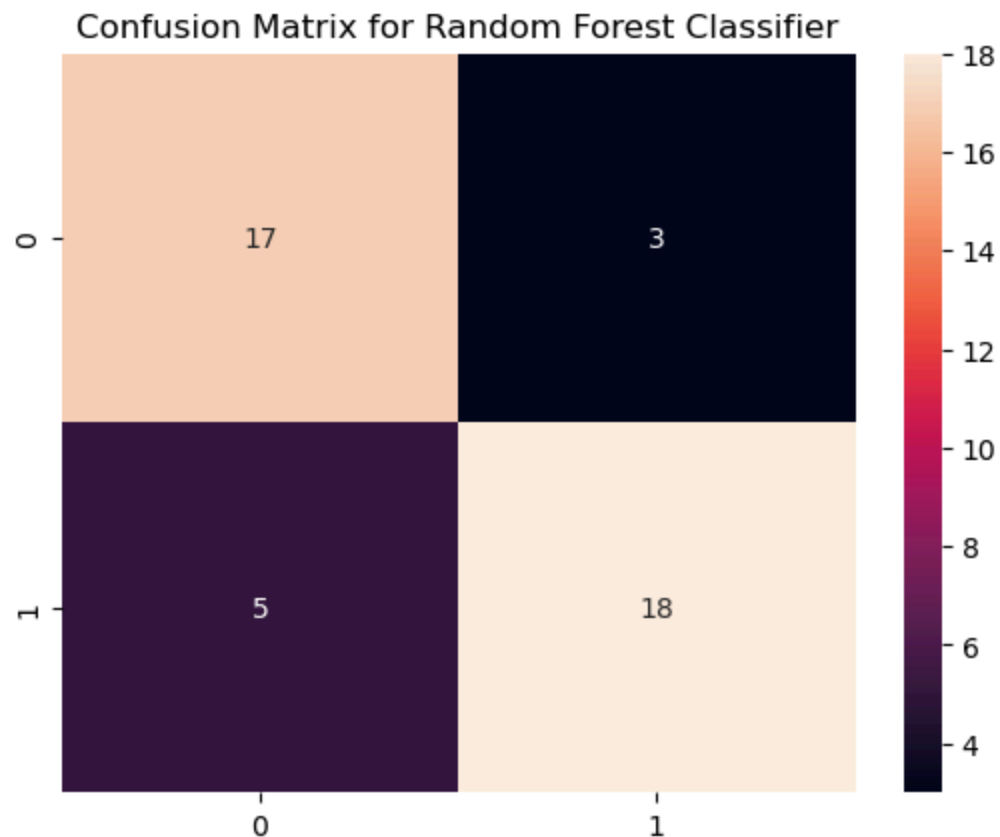
```
In [83]: RFC = RandomForestClassifier(max_depth=2, n_estimators=50, max_features='sqrt')

RFC.fit(X_train, y_train)
rfc_y_pred = RFC.predict(X_test)

rfc_cfm = confusion_matrix(y_test, rfc_y_pred)
print(f"Accuracy of RandomForest is: {accuracy_score(y_test, rfc_y_pred)}")

sns.heatmap(rfc_cfm, annot=True)
plt.title('Confusion Matrix for Random Forest Classifier')
plt.show()
```

Accuracy of RandomForest is: 0.813953488372093



(e) Adaboost with hyperparameter tuning


```
In [97]: from sklearn.ensemble import AdaBoostClassifier
```

```
In [98]: ADB = AdaBoostClassifier(n_estimators = 50, learning_rate=0.01, random_state=0)
ADB.fit(X_train, y_train)
adb_y_pred = ADB.predict(X_test)

adb_cfm = confusion_matrix(y_test, adb_y_pred)
print(f"Accuracy of AdaBoost is: {accuracy_score(y_test, adb_y_pred)}")

sns.heatmap(adb_cfm, annot=True)
plt.title('Confusion Matrix for AdaBoost')
plt.show()
```

Accuracy of AdaBoost is: 0.7906976744186046




(f) Gradientboosting with hyperparameter tuning

```
In [92]: from sklearn.ensemble import GradientBoostingClassifier
gbr_base_model = GradientBoostingClassifier(random_state=0)

gbr_param_grid = {
    'n_estimators': [10, 50, 80, 100, 150, 200, 250, 300],
    'max_depth': [1, 2, 3, 4, 10, 15, 20],
    'learning_rate': [0.01, 0.02, 0.1, 0.05, 0.5, 0.07, 0.9]
}

gbr_gs = GridSearchCV(estimator = gbr_base_model, param_grid = gbr_param_grid, cv=5, n_jobs = -1)
gbr_gs.fit(X_train, y_train)
```


Out[92]:



```
In [93]: print(f'Best parameters for "Gradient Boosting Classifier": {gbr_gs.best_params_}')
print(f'Best Accuracy: {gbr_gs.best_score_}')
```

Best parameters for "Gradient Boosting Classifier": {'learning_rate': 0.01, 'max_depth': 1, 'n_estimators': 300}
 Best Accuracy: 0.7752228163992869

```
In [94]: GBR = GradientBoostingClassifier(learning_rate=0.01, n_estimators=300, max_depth = 1, random_state=0)
GBR.fit(X_train, y_train)
```

Out[94]:

```

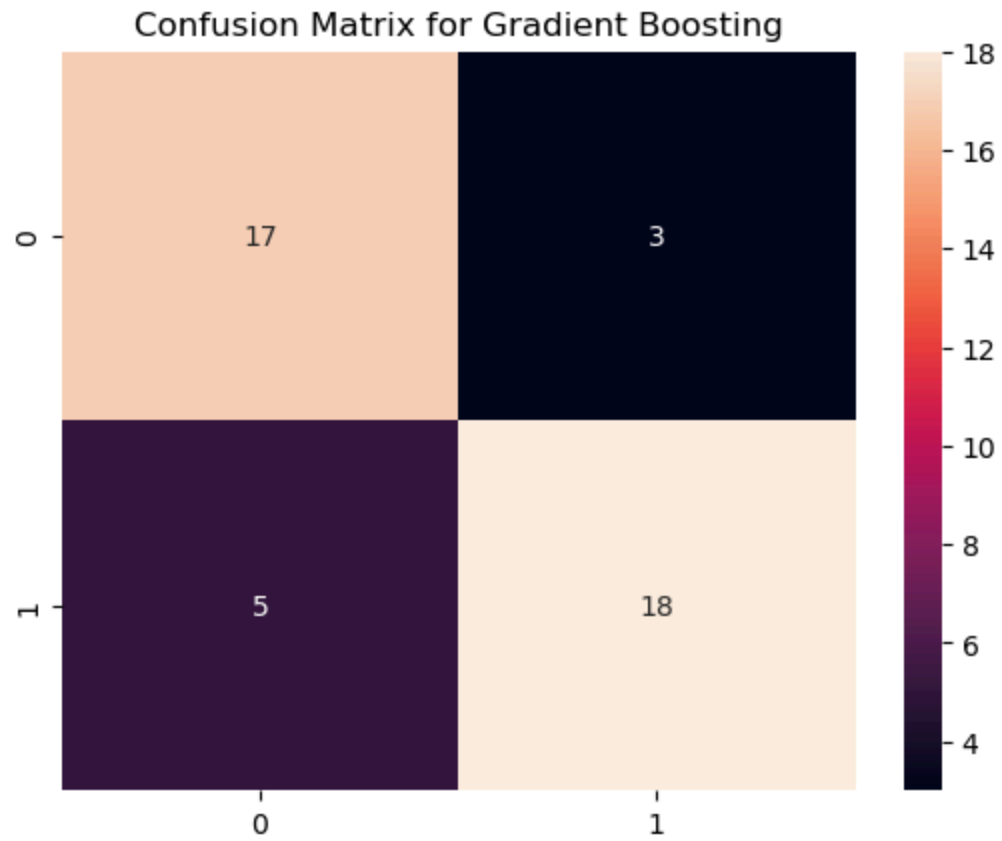
▼ GradientBoostingClassifier
GradientBoostingClassifier(learning_rate=0.01, max_depth=1, n_estimators=300,
                           random_state=0)
  
```

```
In [95]: gbr_y_pred = GBR.predict(X_test)

gbr_cfm = confusion_matrix(y_test, gbr_y_pred)
print(f'Accuracy of Gradient Boosting is: {accuracy_score(y_test, gbr_y_pred)}')

sns.heatmap(gbr_cfm, annot=True)
plt.title('Confusion Matrix for Gradient Boosting')
plt.show()
```

Accuracy of Gradient Boosting is: 0.813953488372093



In []: